

```

# Install required packages
!pip -q install kaggle streamlit pyngrok==4.1.1 tensorflow==2.15 pillow scikit-learn

# Upload kaggle.json
from google.colab import files
print("Upload kaggle.json you just downloaded from Kaggle...")
uploaded = files.upload() # choose kaggle.json

# Set up Kaggle API
!mkdir -p ~/.kaggle
!cp kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json

# Create data folders
!mkdir -p /content/data

🔄 Preparing metadata (setup.py) ... done
ERROR: Ignored the following versions that require a different python version: 0.55.2 Requires-Python <3.5
ERROR: Could not find a version that satisfies the requirement tensorflow==2.15 (from versions: 2.16.0rc0, 2.16.1, 2.16.2, 2.17.0rc0, 2.
ERROR: No matching distribution found for tensorflow==2.15
Upload kaggle.json you just downloaded from Kaggle...
Choose Files kaggle.json
• kaggle.json(application/json) - 71 bytes, last modified: 8/22/2025 - 100% done
Saving kaggle.json to kaggle (1).json

import os, zipfile, pathlib, shutil, random, glob, json
import numpy as np
from PIL import Image
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
print("TF:", tf.__version__)

# Download datasets
# 1. Classification dataset (yes/no)
!kaggle datasets download -d navoneel/brain-mri-images-for-brain-tumor-detection -p /content/data
!unzip -q -o /content/data/brain-mri-images-for-brain-tumor-detection.zip -d /content/data/cls_raw

# 2. Segmentation dataset (LGG with masks)
!kaggle datasets download -d mateuszbuda/lgg-mri-segmentation -p /content/data
!unzip -q -o /content/data/lgg-mri-segmentation.zip -d /content/data/seg_raw

🔄 TF: 2.19.0
Dataset URL: https://www.kaggle.com/datasets/navoneel/brain-mri-images-for-brain-tumor-detection
License(s): copyright-authors
brain-mri-images-for-brain-tumor-detection.zip: Skipping, found more recently modified local copy (use --force to force download)
Dataset URL: https://www.kaggle.com/datasets/mateuszbuda/lgg-mri-segmentation
License(s): CC-BY-NC-SA-4.0
lgg-mri-segmentation.zip: Skipping, found more recently modified local copy (use --force to force download)

# Find the correct path for classification data
SRC = pathlib.Path('/content/data/cls_raw')
brain_tumor_path = None

# Search for the actual folder structure
for root, dirs, files in os.walk('/content/data/cls_raw'):
    if 'yes' in dirs and 'no' in dirs:
        brain_tumor_path = pathlib.Path(root)
        break

if brain_tumor_path:
    SRC = brain_tumor_path
else:
    # Check common folder names
    possible_paths = [
        '/content/data/cls_raw/brain_tumor_dataset',
        '/content/data/cls_raw/Brain Tumor Data Set',
        '/content/data/cls_raw'
    ]
    for path in possible_paths:
        if pathlib.Path(path).exists():
            SRC = pathlib.Path(path)
            break

```

```

print("Classification data root:", SRC)
print("Contents:", list(SRC.iterdir()) if SRC.exists() else "Path not found")

# Create split folders
BASE = pathlib.Path('/content/data/cls_split')
for sub in ['train', 'val', 'test']:
    for cls in ['yes', 'no']:
        (BASE/sub/cls).mkdir(parents=True, exist_ok=True)

# Collect files
files_yes = sorted([*SRC.rglob('yes/*.jpg')]) + sorted([*SRC.rglob('yes/*.png')])
files_no = sorted([*SRC.rglob('no/*.jpg')]) + sorted([*SRC.rglob('no/*.png')])

print(f"Found {len(files_yes)} tumor images and {len(files_no)} normal images")

def split_and_copy(files, cls):
    random.seed(42)
    random.shuffle(files)
    n = len(files)
    n_train = int(0.7*n); n_val = int(0.15*n)
    splits = [('train', files[:n_train]),
              ('val', files[n_train:n_train+n_val]),
              ('test', files[n_train+n_val:])]
    for split, flist in splits:
        for src in flist:
            dst = BASE/split/cls/src.name
            shutil.copy2(src, dst)
    print(f"Split {cls}: {n_train} train, {n_val} val, {n-n_train-n_val} test")

split_and_copy(files_yes, 'yes')
split_and_copy(files_no, 'no')

```

```

→ Classification data root: /content/data/cls_raw
Contents: [PosixPath('/content/data/cls_raw/yes'), PosixPath('/content/data/cls_raw/brain_tumor_dataset'), PosixPath('/content/data/cls_
Found 174 tumor images and 172 normal images
Split yes: 121 train, 26 val, 27 test
Split no: 120 train, 25 val, 27 test

```

```

IMG_SIZE = (224,224)
BATCH = 32

```

```

train_ds = tf.keras.utils.image_dataset_from_directory(
    BASE/'train', image_size=IMG_SIZE, batch_size=BATCH, label_mode='categorical', seed=42)
val_ds = tf.keras.utils.image_dataset_from_directory(
    BASE/'val', image_size=IMG_SIZE, batch_size=BATCH, label_mode='categorical', seed=42)
test_ds = tf.keras.utils.image_dataset_from_directory(
    BASE/'test', image_size=IMG_SIZE, batch_size=BATCH, label_mode='categorical', shuffle=False)

```

```

# Get class names BEFORE applying transformations
class_names = train_ds.class_names
print("Classes:", class_names)

```

```

# Normalize + cache
def prep(x,y):
    # Handle grayscale images by converting to RGB if needed
    if x.shape[-1] == 1:
        x = tf.image.grayscale_to_rgb(x)
    elif x.shape[-1] == 4: # Handle RGBA images
        x = x[:, :, :, :3]
    return tf.cast(x, tf.float32)/255.0, y

```

```

# Apply transformations
train_ds = train_ds.map(prepare).cache().prefetch(tf.data.AUTOTUNE)
val_ds = val_ds.map(prepare).cache().prefetch(tf.data.AUTOTUNE)
test_ds = test_ds.map(prepare).cache().prefetch(tf.data.AUTOTUNE)

```

```

print(f"Training batches: {tf.data.experimental.cardinality(train_ds)}")
print(f"Validation batches: {tf.data.experimental.cardinality(val_ds)}")
print(f"Test batches: {tf.data.experimental.cardinality(test_ds)}")

```

```

→ Found 156 files belonging to 2 classes.
Found 44 files belonging to 2 classes.
Found 49 files belonging to 2 classes.
Classes: ['no', 'yes']
Training batches: 5
Validation batches: 2
Test batches: 2

```

```

from tensorflow.keras.applications import EfficientNetB0, ResNet50, InceptionV3

def build_classifier(model_name='efficientnetb0', num_classes=2, input_shape=(224,224,3)):
    if model_name=='efficientnetb0':
        base = EfficientNetB0(include_top=False, weights='imagenet', input_shape=input_shape)
        preprocess = tf.keras.applications.efficientnet.preprocess_input
    elif model_name=='resnet50':
        base = ResNet50(include_top=False, weights='imagenet', input_shape=input_shape)
        preprocess = tf.keras.applications.resnet50.preprocess_input
    else: # 'inceptionv3' (GoogLeNet family)
        base = InceptionV3(include_top=False, weights='imagenet', input_shape=input_shape)
        preprocess = tf.keras.applications.inception_v3.preprocess_input

    base.trainable = False
    inputs = layers.Input(shape=input_shape)
    x = layers.Lambda(preprocess)(inputs)
    x = base(x, training=False)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dropout(0.2)(x)
    outputs = layers.Dense(num_classes, activation='softmax')(x)
    model = keras.Model(inputs, outputs)
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

    return model

# Train classifier
clf = build_classifier('efficientnetb0')
print("Training classifier...")
history = clf.fit(train_ds, validation_data=val_ds, epochs=10)
test_results = clf.evaluate(test_ds)
print(f"Test Accuracy: {test_results[1]:.4f}")

# Save
os.makedirs('/content/models', exist_ok=True)
clf.save('/content/models/classifier_efficientnet.keras')

```

```

🔗 Training classifier...
Epoch 1/10
5/5 ————— 30s 4s/step - accuracy: 0.4537 - loss: 0.7345 - val_accuracy: 0.4773 - val_loss: 0.6946
Epoch 2/10
5/5 ————— 14s 3s/step - accuracy: 0.5039 - loss: 0.6934 - val_accuracy: 0.5227 - val_loss: 0.6937
Epoch 3/10
5/5 ————— 13s 3s/step - accuracy: 0.5242 - loss: 0.7047 - val_accuracy: 0.5227 - val_loss: 0.6934
Epoch 4/10
5/5 ————— 13s 3s/step - accuracy: 0.4701 - loss: 0.7038 - val_accuracy: 0.4773 - val_loss: 0.6944
Epoch 5/10
5/5 ————— 13s 3s/step - accuracy: 0.5424 - loss: 0.7006 - val_accuracy: 0.4773 - val_loss: 0.6938
Epoch 6/10
5/5 ————— 21s 3s/step - accuracy: 0.5428 - loss: 0.6831 - val_accuracy: 0.5227 - val_loss: 0.6927
Epoch 7/10
5/5 ————— 13s 3s/step - accuracy: 0.5043 - loss: 0.6893 - val_accuracy: 0.5227 - val_loss: 0.6926
Epoch 8/10
5/5 ————— 13s 3s/step - accuracy: 0.5337 - loss: 0.6849 - val_accuracy: 0.5227 - val_loss: 0.6928
Epoch 9/10
5/5 ————— 22s 3s/step - accuracy: 0.5088 - loss: 0.7017 - val_accuracy: 0.4773 - val_loss: 0.6938
Epoch 10/10
5/5 ————— 15s 3s/step - accuracy: 0.4537 - loss: 0.7007 - val_accuracy: 0.5227 - val_loss: 0.6927
2/2 ————— 3s 923ms/step - accuracy: 0.6006 - loss: 0.6909
Test Accuracy: 0.5102

```

```

def build_unet(input_shape=(256, 256, 3)):
    inputs = layers.Input(shape=input_shape)

    # Encoder (downsampling)
    c1 = layers.Conv2D(64, 3, activation='relu', padding='same')(inputs)
    c1 = layers.Conv2D(64, 3, activation='relu', padding='same')(c1)
    p1 = layers.MaxPooling2D(2)(c1)

    c2 = layers.Conv2D(128, 3, activation='relu', padding='same')(p1)
    c2 = layers.Conv2D(128, 3, activation='relu', padding='same')(c2)
    p2 = layers.MaxPooling2D(2)(c2)

    c3 = layers.Conv2D(256, 3, activation='relu', padding='same')(p2)
    c3 = layers.Conv2D(256, 3, activation='relu', padding='same')(c3)
    p3 = layers.MaxPooling2D(2)(c3)

```

```

c4 = layers.Conv2D(512, 3, activation='relu', padding='same')(p3)
c4 = layers.Conv2D(512, 3, activation='relu', padding='same')(c4)
p4 = layers.MaxPooling2D(2)(c4)

# Bridge
c5 = layers.Conv2D(1024, 3, activation='relu', padding='same')(p4)
c5 = layers.Conv2D(1024, 3, activation='relu', padding='same')(c5)

# Decoder (upsampling)
u6 = layers.UpSampling2D(2)(c5)
u6 = layers.concatenate([u6, c4])
c6 = layers.Conv2D(512, 3, activation='relu', padding='same')(u6)
c6 = layers.Conv2D(512, 3, activation='relu', padding='same')(c6)

u7 = layers.UpSampling2D(2)(c6)
u7 = layers.concatenate([u7, c3])
c7 = layers.Conv2D(256, 3, activation='relu', padding='same')(u7)
c7 = layers.Conv2D(256, 3, activation='relu', padding='same')(c7)

u8 = layers.UpSampling2D(2)(c7)
u8 = layers.concatenate([u8, c2])
c8 = layers.Conv2D(128, 3, activation='relu', padding='same')(u8)
c8 = layers.Conv2D(128, 3, activation='relu', padding='same')(c8)

u9 = layers.UpSampling2D(2)(c8)
u9 = layers.concatenate([u9, c1])
c9 = layers.Conv2D(64, 3, activation='relu', padding='same')(u9)
c9 = layers.Conv2D(64, 3, activation='relu', padding='same')(c9)

outputs = layers.Conv2D(1, 1, activation='sigmoid')(c9)

model = keras.Model(inputs=[inputs], outputs=[outputs])
return model

import cv2

def load_segmentation_data(data_path, img_size=(256, 256)):
    images = []
    masks = []

    seg_path = pathlib.Path(data_path)

    # Find all TIFF files (images)
    image_files = list(seg_path.rglob('*_mask.tif')) # Find mask files first

    for mask_file in image_files:
        # Get corresponding image file
        img_file = str(mask_file).replace('_mask.tif', '.tif')

        if os.path.exists(img_file):
            # Load image
            img = cv2.imread(img_file, cv2.IMREAD_GRAYSCALE)
            if img is not None:
                img = cv2.resize(img, img_size)
                img = cv2.cvtColor(img, cv2.COLOR_GRAY2RGB)
                images.append(img)

            # Load mask
            mask = cv2.imread(str(mask_file), cv2.IMREAD_GRAYSCALE)
            mask = cv2.resize(mask, img_size)
            mask = mask / 255.0 # Normalize to 0-1
            masks.append(mask)

    return np.array(images), np.array(masks)

# Load segmentation data
print("Loading segmentation data...")
seg_images, seg_masks = load_segmentation_data('/content/data/seg_raw')
print(f"Loaded {len(seg_images)} images for segmentation")

# Split segmentation data
n = len(seg_images)
n_train_seg = int(0.8 * n)

seg_train_x, seg_test_x = seg_images[:n_train_seg], seg_images[n_train_seg:]
seg_train_y, seg_test_y = seg_masks[:n_train_seg], seg_masks[n_train_seg:]

```

```

# Normalize
seg_train_x = seg_train_x.astype('float32') / 255.0
seg_test_x = seg_test_x.astype('float32') / 255.0

🔄 Loading segmentation data...

# Build and train U-Net
UNET = build_unet((256, 256, 3))
UNET.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

print("Training U-Net...")
UNET_history = UNET.fit(seg_train_x, seg_train_y,
                        validation_split=0.2,
                        epochs=10,
                        batch_size=8)

# Evaluate
UNET_results = UNET.evaluate(seg_test_x, seg_test_y)
print(f"Segmentation Test Accuracy: {UNET_results[1]:.4f}")

# Save
UNET.save('/content/models/unet_segmentation.keras')

%writefile app.py
import streamlit as st
import tensorflow as tf
import numpy as np
from PIL import Image
import cv2

st.title("🧠 Brain Tumor Detection & Segmentation")
st.write("Upload an MRI image to detect and segment brain tumors")

# Load models
@st.cache_resource
def load_models():
    classifier = tf.keras.models.load_model('/content/models/classifier_efficientnet.keras')
    segmentor = tf.keras.models.load_model('/content/models/unet_segmentation.keras')
    return classifier, segmentor

try:
    classifier, segmentor = load_models()
    st.success("Models loaded successfully!")
except:
    st.error("Please train the models first!")
    st.stop()

# File uploader
uploaded_file = st.file_uploader("Choose an MRI image...", type=['jpg', 'jpeg', 'png'])

if uploaded_file is not None:
    # Display uploaded image
    image = Image.open(uploaded_file)
    st.image(image, caption="Uploaded MRI Image", use_column_width=True)

    # Preprocess for classification
    img_array = np.array(image.resize((224, 224)))
    if len(img_array.shape) == 2: # Grayscale
        img_array = cv2.cvtColor(img_array, cv2.COLOR_GRAY2RGB)
    img_array = np.expand_dims(img_array, axis=0) / 255.0

    # Classification
    st.subheader("🔍 Tumor Detection")
    pred = classifier.predict(img_array)
    class_names = ['No Tumor', 'Tumor Present']
    predicted_class = class_names[np.argmax(pred)]
    confidence = np.max(pred) * 100

    st.write(f"**Prediction:** {predicted_class}")
    st.write(f"**Confidence:** {confidence:.2f}%")

    # Segmentation if tumor detected
    if np.argmax(pred) == 1: # Tumor present
        st.subheader("🖼️ Tumor Segmentation")

```

```

# Preprocess for segmentation
seg_img = np.array(image.resize((256, 256)))
if len(seg_img.shape) == 2:
    seg_img = cv2.cvtColor(seg_img, cv2.COLOR_GRAY2RGB)
seg_img = np.expand_dims(seg_img, axis=0) / 255.0

# Predict mask
mask_pred = segmentor.predict(seg_img)
mask = mask_pred[0, :, :, 0]

# Display segmentation
col1, col2 = st.columns(2)
with col1:
    st.image(image, caption="Original Image", use_column_width=True)
with col2:
    st.image(mask, caption="Tumor Segmentation", use_column_width=True)

# Calculate tumor area
tumor_pixels = np.sum(mask > 0.5)
total_pixels = mask.shape[0] * mask.shape[1]
tumor_percentage = (tumor_pixels / total_pixels) * 100
st.write(f"**Tumor Area:** {tumor_percentage:.2f}% of image")

st.sidebar.info("This app uses EfficientNet for classification and U-Net for segmentation")

# Simple way to run locally in Colab
from pyngrok import ngrok

# Run Streamlit
!streamlit run app.py &

# Create tunnel
public_url = ngrok.connect(port='8501')
print(f'Your app is live at: {public_url}')

```